

```
In [1]: ▶ import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import backend as K
from tensorflow.keras.layers import Dense, Activation, Dropout, Conv2D, MaxPool
from tensorflow.keras.optimizers import Adam, Adamax
from tensorflow.keras.metrics import categorical_crossentropy
from tensorflow.keras import regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Model, load_model, Sequential
import numpy as np
import pandas as pd
import shutil
import time
import cv2 as cv2
from tqdm import tqdm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from matplotlib.pyplot import imshow
import seaborn as sns
sns.set_style('darkgrid')
from PIL import Image
from sklearn.metrics import confusion_matrix, classification_report
from IPython.core.display import display, HTML
# stop annoying tensorflow warning messages
import logging
logging.getLogger("tensorflow").setLevel(logging.ERROR)
print ('modules loaded')
```

modules loaded

```
C:\Users\mahir\AppData\Local\Temp\ipykernel_10900\1633298320.py:26: DeprecationWarning: Importing display from IPython.core.display is deprecated since IPython 7.14, please import from IPython display
  from IPython.core.display import display, HTML
```

```
In [2]: ▶ def show_image_samples(gen):
    t_dict=gen.class_indices
    classes=list(t_dict.keys())
    images,labels=next(gen) # get a sample batch from the generator
    plt.figure(figsize=(20, 20))
    length=len(labels)
    if length<25: #show maximum of 25 images
        r=length
    else:
        r=25
    for i in range(r):
        plt.subplot(5, 5, i + 1)
        image=images[i]/255
        plt.imshow(image)
        index=np.argmax(labels[i])
        class_name=classes[index]
        plt.title(class_name, color='blue', fontsize=12)
        plt.axis('off')
    plt.show()
```

```
In [3]: ▶ def show_images(tdir):
    classlist=os.listdir(tdir)
    length=len(classlist)
    columns=5
    rows=int(np.ceil(length/columns))
    plt.figure(figsize=(20, rows * 4))
    for i, class in enumerate(classlist):
        classpath=os.path.join(tdir, class)
        imgpath=os.path.join(classpath, '1.jpg')
        img=plt.imread(imgpath)
        plt.subplot(rows, columns, i+1)
        plt.axis('off')
        plt.title(class, color='blue', fontsize=12)
        plt.imshow(img)
```

```
In [4]: ▶ def print_in_color(txt_msg,fore_tuple,back_tuple,):
    #prints the text_msg in the foreground color specified by fore_tuple with
    #text_msg is the text, fore_tuple is foreground color tuple (r,g,b), back
    rf,gf,bf=fore_tuple
    rb,gb,bb=back_tuple
    msg='{0}' + txt_msg
    mat='\33[38;2;' + str(rf) + ';' + str(gf) + ';' + str(bf) + ';48;2;' + str
    print(msg .format(mat), flush=True)
    print('\33[0m', flush=True) # returns default print color to back to black
    return
```

```

In [5]: class LRA(keras.callbacks.Callback):
    def __init__(self,model, base_model, patience,stop_patience, threshold, f
        super(LRA, self).__init__()
        self.model=model
        self.base_model=base_model
        self.patience=patience # specifies how many epochs without improvemen
        self.stop_patience=stop_patience # specifies how many times to adjust
        self.threshold=threshold # specifies training accuracy threshold when
        self.factor=factor # factor by which to reduce the learning rate
        self.dwell=dwell
        self.batches=batches # number of training batch to runn per epoch
        self.initial_epoch=initial_epoch
        self.epochs=epochs
        self.ask_epoch=ask_epoch
        self.ask_epoch_initial=ask_epoch # save this value to restore if rest
        # callback variables
        self.count=0 # how many times lr has been reduced without improvement
        self.stop_count=0
        self.best_epoch=1 # epoch with the lowest loss
        self.initial_lr=float(tf.keras.backend.get_value(model.optimizer.lr))
        self.highest_tracc=0.0 # set highest training accuracy to 0 initially
        self.lowest_vloss=np.inf # set lowest validation loss to infinity ini
        self.best_weights=self.model.get_weights() # set best weights to mode
        self.initial_weights=self.model.get_weights() # save initial weight

    def on_train_begin(self, logs=None):
        if self.base_model != None:
            status=base_model.trainable
            if status:
                msg='initializing callback starting training with base_model
            else:
                msg='initializing callback starting training with base_model
        else:
            msg='initialing callback and starting training'
        print_in_color (msg, (244, 252, 3), (55,65,80))
        msg='{0:^8s}{1:^10s}{2:^9s}{3:^9s}{4:^9s}{5:^9s}{6:^9s}{7:^10s}{8:10s

        print_in_color(msg, (244,252,3), (55,65,80))
        self.start_time= time.time()

    def on_train_end(self, logs=None):
        stop_time=time.time()
        tr_duration= stop_time- self.start_time
        hours = tr_duration // 3600
        minutes = (tr_duration - (hours * 3600)) // 60
        seconds = tr_duration - ((hours * 3600) + (minutes * 60))

        self.model.set_weights(self.best_weights) # set the weights of the mo
        msg=f'Training is completed - model is set with weights from epoch {s
        print_in_color(msg, (0,255,0), (55,65,80))
        msg = f'training elapsed time was {str(hours)} hours, {minutes:4.1f}
        print_in_color(msg, (0,255,0), (55,65,80))

    def on_train_batch_end(self, batch, logs=None):
        acc=logs.get('accuracy')* 100 # get training accuracy
        loss=logs.get('loss')

```

```

msg='{0:20s}processing batch {1:4s} of {2:5s} accuracy= {3:8.3f}  los
print(msg, '\r', end='') # prints over on the same line to show runni

def on_epoch_begin(self,epoch, logs=None):
    self.now= time.time()

def on_epoch_end(self, epoch, logs=None): # method runs on the end of ea
    later=time.time()
    duration=later-self.now
    lr=float(tf.keras.backend.get_value(self.model.optimizer.lr)) # get t
    current_lr=lr
    v_loss=logs.get('val_loss') # get the validation loss for this epoch
    acc=logs.get('accuracy') # get training accuracy
    v_acc=logs.get('val_accuracy')
    loss=logs.get('loss')
    if acc < self.threshold: # if training accuracy is below threshold ad
        monitor='accuracy'
        if epoch ==0:
            pimprov=0.0
        else:
            pimprov= (acc-self.highest_tracc )*100/self.highest_tracc
        if acc>self.highest_tracc: # training accuracy improved in the ep
            self.highest_tracc=acc # set new highest training accuracy
            self.best_weights=self.model.get_weights() # traing accuracy
            self.count=0 # set count to 0 since training accuracy improve
            self.stop_count=0 # set stop counter to 0
            if v_loss<self.lowest_vloss:
                self.lowest_vloss=v_loss
            color= (0,255,0)
            self.best_epoch=epoch + 1 # set the value of best epoch for
        else:
            # training accuracy did not improve check if this has happene
            # if so adjust Learning rate
            if self.count>=self.patience -1: # Lr should be adjusted
                color=(245, 170, 66)
                lr= lr* self.factor # adjust the Learning by factor
                tf.keras.backend.set_value(self.model.optimizer.lr, lr) #
                self.count=0 # reset the count to 0
                self.stop_count=self.stop_count + 1 # count the number of
                self.count=0 # reset counter
                if self.dwell:
                    self.model.set_weights(self.best_weights) # return to
                else:
                    if v_loss<self.lowest_vloss:
                        self.lowest_vloss=v_loss
            else:
                self.count=self.count +1 # increment patience counter
        else: # training accuracy is above threshold so adjust Learning rate
            monitor='val_loss'
            if epoch ==0:
                pimprov=0.0
            else:
                pimprov= (self.lowest_vloss- v_loss )*100/self.lowest_vloss
            if v_loss< self.lowest_vloss: # check if the validation loss impr
                self.lowest_vloss=v_loss # replace lowest validation loss wit
                self.best_weights=self.model.get_weights() # validation loss
                self.count=0 # reset count since validation loss improved

```

```

self.stop_count=0
color=(0,255,0)
self.best_epoch=epoch + 1 # set the value of the best epoch to
else: # validation loss did not improve
    if self.count>=self.patience-1: # need to adjust lr
        color=(245, 170, 66)
        lr=lr * self.factor # adjust the Learning rate
        self.stop_count=self.stop_count + 1 # increment stop count
        self.count=0 # reset counter
        tf.keras.backend.set_value(self.model.optimizer.lr, lr) #
        if self.dwell:
            self.model.set_weights(self.best_weights) # return to
    else:
        self.count =self.count +1 # increment the patience counter
    if acc>self.highest_tracc:
        self.highest_tracc= acc
msg=f'{str(epoch+1):^3s}/{str(self.epochs):4s} {loss:^9.3f}{acc*100:^9.3f}'
print_in_color (msg,color, (55,65,80))
if self.stop_count> self.stop_patience - 1: # check if Learning rate
    msg=f' training has been halted at epoch {epoch + 1} after {self.
    print_in_color(msg, (0,255,255), (55,65,80))
    self.model.stop_training = True # stop training
else:
    if self.ask_epoch !=None:
        if epoch + 1 >= self.ask_epoch:
            if base_model.trainable:
                msg='enter H to halt training or an integer for number
            else:
                msg='enter H to halt training ,F to fine tune model,
                print_in_color(msg, (0,255,255), (55,65,80))
                ans=input('')
                if ans=='H' or ans=='h':
                    msg=f'training has been halted at epoch {epoch + 1} c
                    print_in_color(msg, (0,255,255), (55,65,80))
                    self.model.stop_training = True # stop training
                elif ans == 'F' or ans=='f':
                    if base_model.trainable:
                        msg='base_model is already set as trainable'
                    else:
                        msg='setting base_model as trainable for fine tun
                        self.base_model.trainable=True
                        print_in_color(msg, (0, 255,255), (55,65,80))
                        msg='{0:^8s}{1:^10s}{2:^9s}{3:^9s}{4:^9s}{5:^9s}{6:^9
                        print_in_color(msg, (244,252,3), (55,65,80))
                        self.count=0
                        self.stop_count=0
                        self.ask_epoch = epoch + 1 + self.ask_epoch_initial
            else:
                ans=int(ans)
                self.ask_epoch +=ans
                msg=f' training will continue until epoch ' + str(sel
                print_in_color(msg, (0, 255,255), (55,65,80))
                msg='{0:^8s}{1:^10s}{2:^9s}{3:^9s}{4:^9s}{5:^9s}{6:^9
                print_in_color(msg, (244,252,3), (55,65,80))

```

```

In [6]: ▶ def tr_plot(tr_data, start_epoch):
    #Plot the training and validation data
    tacc=tr_data.history['accuracy']
    tloss=tr_data.history['loss']
    vacc=tr_data.history['val_accuracy']
    vloss=tr_data.history['val_loss']
    Epoch_count=len(tacc)+ start_epoch
    Epochs=[]
    for i in range (start_epoch ,Epoch_count):
        Epochs.append(i+1)
    index_loss=np.argmin(vloss)# this is the epoch with the lowest validation loss
    val_lowest=vloss[index_loss]
    index_acc=np.argmax(vacc)
    acc_highest=vacc[index_acc]
    plt.style.use('fivethirtyeight')
    sc_label='best epoch= ' + str(index_loss+1 +start_epoch)
    vc_label='best epoch= ' + str(index_acc + 1+ start_epoch)
    fig,axes=plt.subplots(nrows=1, ncols=2, figsize=(20,8))
    axes[0].plot(Epochs,tloss, 'r', label='Training loss')
    axes[0].plot(Epochs,vloss,'g',label='Validation loss' )
    axes[0].scatter(index_loss+1 +start_epoch,val_lowest, s=150, c= 'blue', 1
    axes[0].set_title('Training and Validation Loss')
    axes[0].set_xlabel('Epochs')
    axes[0].set_ylabel('Loss')
    axes[0].legend()
    axes[1].plot (Epochs,tacc,'r',label= 'Training Accuracy')
    axes[1].plot (Epochs,vacc,'g',label= 'Validation Accuracy')
    axes[1].scatter(index_acc+1 +start_epoch,acc_highest, s=150, c= 'blue', 1
    axes[1].set_title('Training and Validation Accuracy')
    axes[1].set_xlabel('Epochs')
    axes[1].set_ylabel('Accuracy')
    axes[1].legend()
    plt.tight_layout
    #plt.style.use('fivethirtyeight')
    plt.show()

```

```

In [7]: ▶ def print_info( test_gen, preds, print_code, save_dir, subject ):
    class_dict=test_gen.class_indices
    labels= test_gen.labels
    file_names= test_gen filenames
    error_list=[]
    true_class=[]
    pred_class=[]
    prob_list=[]
    new_dict={}
    error_indices=[]
    y_pred=[]
    for key,value in class_dict.items():
        new_dict[value]=key           # dictionary {integer of class number
    # store new_dict as a text file in the save_dir
    classes=list(new_dict.values())    # list of string of class names
    errors=0
    for i, p in enumerate(preds):
        pred_index=np.argmax(p)
        true_index=labels[i] # Labels are integer values
        if pred_index != true_index: # a misclassification has occurred
            error_list.append(file_names[i])
            true_class.append(new_dict[true_index])
            pred_class.append(new_dict[pred_index])
            prob_list.append(p[pred_index])
            error_indices.append(true_index)
            errors=errors + 1
        y_pred.append(pred_index)
    tests=len(preds)
    acc= (1-errors/tests) *100
    msg= f'There were {errors} errors in {tests} test cases Model accuracy= {
    print_in_color(msg,(0,255,255),(55,65,80))
    if print_code !=0:
        if errors>0:
            if print_code>errors:
                r=errors
            else:
                r=print_code
            msg='{0:^28s}{1:^28s}{2:^28s}{3:^16s}'.format('Filename', 'Predicted', 'True', 'Prob')
            print_in_color(msg, (0,255,0),(55,65,80))
            for i in range(r):
                split1=os.path.split(error_list[i])
                split2=os.path.split(split1[0])
                fname=split2[1] + '/' + split1[1]
                msg='{0:^28s}{1:^28s}{2:^28s}{3:4s}{4:^6.4f}'.format(fname, p[i], true_class[i], prob_list[i])
                print_in_color(msg, (255,255,255), (55,65,60))
                #print(error_list[i] , pred_class[i], true_class[i], prob_list[i])
            else:
                msg='With accuracy of 100 % there are no errors to print'
                print_in_color(msg, (0,255,0),(55,65,80))
    if errors>0:
        plot_bar=[]
        plot_class=[]
        for key, value in new_dict.items():
            count=error_indices.count(key)
            if count!=0:
                plot_bar.append(count) # List containing how many times a class

```

```

        plot_class.append(value)    # stores the class
fig=plt.figure()
fig.set_figheight(len(plot_class)/3)
fig.set_figwidth(10)
plt.style.use('fivethirtyeight')
for i in range(0, len(plot_class)):
    c=plot_class[i]
    x=plot_bar[i]
    plt.barh(c, x, )
    plt.title( ' Errors by Class on Test Set')
y_true= np.array(labels)
y_pred=np.array(y_pred)
if len(classes)<= 30:
    # create a confusion matrix
    cm = confusion_matrix(y_true, y_pred )
    length=len(classes)
    if length<8:
        fig_width=8
        fig_height=8
    else:
        fig_width= int(length * .5)
        fig_height= int(length * .5)
    plt.figure(figsize=(fig_width, fig_height))
    sns.heatmap(cm, annot=True, vmin=0, fmt='g', cmap='Blues', cbar=False)
    plt.xticks(np.arange(length)+.5, classes, rotation= 90)
    plt.yticks(np.arange(length)+.5, classes, rotation=0)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title("Confusion Matrix")
    plt.show()
clr = classification_report(y_true, y_pred, target_names=classes, digits=2)
print("Classification Report:\n-----\n", clr)
return acc/100

```



```

In [8]: ▶ def saver(save_path, model, model_name, subject, accuracy, img_size, scalar, g
# first save the model
save_id=str(model_name + '-' + subject + '-' + str(acc)[:str(acc).rfind('
model_save_loc=os.path.join(save_path, save_id)
model.save(model_save_loc)
print_in_color('model was saved as ' + model_save_loc, (0,255,0),(55,65,
# now create the class_df and convert to csv file
class_dict=generator.class_indices
height=[]
width=[]
scale=[]
for i in range(len(class_dict)):
    height.append(img_size[0])
    width.append(img_size[1])
    scale.append(scalar)
Index_series=pd.Series(list(class_dict.values()), name='class_index')
Class_series=pd.Series(list(class_dict.keys()), name='class')
Height_series=pd.Series(height, name='height')
Width_series=pd.Series(width, name='width')
Scale_series=pd.Series(scale, name='scale by')
class_df=pd.concat([Index_series, Class_series, Height_series, Width_seri
csv_name='class_dict.csv'
csv_save_loc=os.path.join(save_path, csv_name)
class_df.to_csv(csv_save_loc, index=False)
print_in_color('class csv file was saved as ' + csv_save_loc, (0,255,0),
return model_save_loc, csv_save_loc

```

```

In [9]: ▶ def predictor(sdir, csv_path, model_path, averaged=True, verbose=True):
    # read in the csv file
    class_df=pd.read_csv(csv_path)
    class_count=len(class_df['class'].unique())
    img_height=int(class_df['height'].iloc[0])
    img_width =int(class_df['width'].iloc[0])
    img_size=(img_width, img_height)
    scale=class_df['scale by'].iloc[0]
    # determine value to scale image pixels by
    try:
        s=int(scale)
        s2=1
        s1=0
    except:
        split=scale.split('-')
        s1=float(split[1])
        s2=float(split[0].split('*')[1])
    path_list=[]
    paths=os.listdir(sdir)
    for f in paths:
        path_list.append(os.path.join(sdir,f))
    if verbose:
        print (' Model is being loaded- this will take about 10 seconds')
    model=load_model(model_path)
    image_count=len(path_list)
    image_list=[]
    file_list=[]
    good_image_count=0
    for i in range (image_count):
        try:
            img=cv2.imread(path_list[i])
            img=cv2.resize(img, img_size)
            img=cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
            good_image_count +=1
            img=img*s2 - s1
            image_list.append(img)
            file_name=os.path.split(path_list[i])[1]
            file_list.append(file_name)
        except:
            if verbose:
                print ( path_list[i], ' is an invalid image file')
    if good_image_count==1: # if only a single image need to expand dimension
        averaged=True
    image_array=np.array(image_list)
    # make predictions on images, sum the probabilities of each class then find
    # highest probability
    preds=model.predict(image_array)
    if averaged:
        psum=[]
        for i in range (class_count): # create all 0 values list
            psum.append(0)
        for p in preds: # iterate over all predictions
            for i in range (class_count):
                psum[i]=psum[i] + p[i] # sum the probabilities
        index=np.argmax(psum) # find the class index with the highest probability
        klass=class_df['class'].iloc[index] # get the class name that corresponds

```

```

prob=psum[index]/good_image_count * 100 # get the probability average
# to show the correct image run predict again and select first image
for img in image_array: #iterate through the images
    test_img=np.expand_dims(img, axis=0) # since it is a single image
    test_index=np.argmax(model.predict(test_img)) # for this image find
    if test_index== index: # see if this image has the same index as
        if verbose: # show image and print result if verbose=1
            plt.axis('off')
            plt.imshow(img) # show the image
            print (f'predicted species is {klass} with a probability
                    {prob}')
            break # found an image that represents the predicted class
    return klass, prob, img, None
else: # create individual predictions for each image
    pred_class=[]
    prob_list=[]
    for i, p in enumerate(preds):
        index=np.argmax(p) # find the class index with the highest probability
        klass=class_df['class'].iloc[index] # get the class name that corresponds
        image_file= file_list[i]
        pred_class.append(klass)
        prob_list.append(p[index])
    Fseries=pd.Series(file_list, name='image file')
    Lseries=pd.Series(pred_class, name='species')
    Pseries=pd.Series(prob_list, name='probability')
    df=pd.concat([Fseries, Lseries, Pseries], axis=1)
    if verbose:
        length= len(df)
        print (df.head(length))
    return None, None, None, df

```

```

In [10]: ▶ def trim (df, max_size, min_size, column):
df=df.copy()
original_class_count= len(list(df[column].unique()))
print ('Original Number of classes in dataframe: ', original_class_count)
sample_list=[]
groups=df.groupby(column)
for label in df[column].unique():
    group=groups.get_group(label)
    sample_count=len(group)
    if sample_count> max_size :
        strat=group[column]
        samples, _=train_test_split(group, train_size=max_size, shuffle=True)
        sample_list.append(samples)
    elif sample_count>= min_size:
        sample_list.append(group)
df=pd.concat(sample_list, axis=0).reset_index(drop=True)
final_class_count= len(list(df[column].unique()))
if final_class_count != original_class_count:
    print ('*** WARNING*** dataframe has a reduced number of classes' )
balance=list(df[column].value_counts())
print (balance)
return df

```

```

In [11]: ▶ def balance(train_df,max_samples, min_samples, column, working_dir, image_size):
    train_df=train_df.copy()
    train_df=trim (train_df, max_samples, min_samples, column)
    # make directories to store augmented images
    aug_dir=os.path.join(working_dir, 'aug')
    if os.path.isdir(aug_dir):
        shutil.rmtree(aug_dir)
    os.mkdir(aug_dir)
    for label in train_df['labels'].unique():
        dir_path=os.path.join(aug_dir,label)
        os.mkdir(dir_path)
    # create and store the augmented images
    total=0
    gen=ImageDataGenerator(horizontal_flip=True, rotation_range=20, width_shift_range=.2, height_shift_range=.2, zoom_range=.2)
    groups=train_df.groupby('labels') # group by class
    for label in train_df['labels'].unique(): # for every class
        group=groups.get_group(label) # a dataframe holding only rows with the given label
        sample_count=len(group) # determine how many samples there are in the class
        if sample_count< max_samples: # if the class has less than target number of samples
            aug_img_count=0
            delta=max_samples-sample_count # number of augmented images to create
            target_dir=os.path.join(aug_dir, label) # define where to write the augmented images
            aug_gen=gen.flow_from_dataframe( group, x_col='filepaths', y_col='labels', class_mode=None, batch_size=1, shuffle=True, save_to_dir=target_dir, save_prefix='', save_format='jpg')

            while aug_img_count<delta:
                images=next(aug_gen)
                aug_img_count += len(images)
            total +=aug_img_count
    print('Total Augmented images created= ', total)
    # create aug_df and merge with train_df to create composite training set
    if total>0:
        aug_fpaths=[]
        aug_labels=[]
        classlist=os.listdir(aug_dir)
        for klass in classlist:
            classpath=os.path.join(aug_dir, klass)
            flist=os.listdir(classpath)
            for f in flist:
                fpath=os.path.join(classpath,f)
                aug_fpaths.append(fpath)
                aug_labels.append(klass)
        Fseries=pd.Series(aug_fpaths, name='filepaths')
        Lseries=pd.Series(aug_labels, name='labels')
        aug_df=pd.concat([Fseries, Lseries], axis=1)
        train_df=pd.concat([train_df,aug_df], axis=0).reset_index(drop=True)


    print (list(train_df['labels'].value_counts()) )
    return train_df

```

```
In [13]: ▶ img_path=r'C:/Users/mahir/Documents/All the dreams in the world/Thesis Project/Rocket Man.png'
img=plt.imread(img_path)
print ('Input image shape is ',img.shape)
plt.axis('off')
imshow(img)
```

Input image shape is (450, 450, 3)

Out[13]: <matplotlib.image.AxesImage at 0x21e21ec9820>



```

In [14]: ▶ def preprocess (sdir, trsplit, vsplit):
    filepaths=[]
    labels=[]
    folds=os.listdir(sdir)
    for fold in folds:
        foldpath=os.path.join(sdir, fold)
        classlist=os.listdir(foldpath)
        for class in classlist:
            classpath=os.path.join(foldpath, class)
            flist=os.listdir(classpath)
            for f in flist:
                fpath=os.path.join(classpath, f)
                filepaths.append(fpath)
                labels.append(class)
    Fseries=pd.Series(filepaths, name='filepaths')
    Lseries=pd.Series(labels, name='labels')
    df=pd.concat([Fseries, Lseries], axis=1)
    dsplit=vsplit/(1-trsplit)
    strat=df['labels']
    train_df, dummy_df=train_test_split(df, train_size=trsplit, shuffle=True,
    strat=dummy_df['labels'])
    valid_df, test_df= train_test_split(dummy_df, train_size=dsplit, shuffle=
    print('train_df length: ', len(train_df), ' test_df length: ', len(test_d
    # check that each dataframe has the same number of classes to prevent mo
    trcount=len(train_df['labels'].unique())
    tecount=len(test_df['labels'].unique())
    vcount=len(valid_df['labels'].unique())
    if trcount < tecount :
        msg='** WARNING ** number of classes in training set is less than the
        print_in_color(msg, (255,0,0), (55,65,80))
        msg='This will throw an error in either model.evaluate or model.predi
        print_in_color(msg, (255,0,0), (55,65,80))
    if trcount != vcount:
        msg='** WARNING ** number of classes in training set not equal to num
        print_in_color(msg, (255,0,0), (55,65,80))
        msg=' this will throw an error in model.fit'
        print_in_color(msg, (255,0,0), (55,65,80))
        print ('train df class count: ', trcount, 'test df class count: ', te
        ans=input('Enter C to continue execution or H to halt execution')
        if ans == 'H' or ans == 'h':
            print_in_color('Halting Execution', (255,0,0), (55,65,80))
            import sys
            sys.exit('program halted by user')
    print(list(train_df['labels'].value_counts()))
    return train_df, test_df, valid_df

```

```

In [17]: ▶ sdir=r'C:/Users/mahir/Documents/All the dreams in the world/Thesis Project/te
    trsplit=.9
    vsplit=.05
    train_df, test_df, valid_df= preprocess(sdir,trsplit, vsplit)

```

```

train_df length: 9594    test_df length: 534    valid_df length: 533
[6544, 3050]

```


```
In [18]: max_samples= 3050
min_samples=0
column='labels'
working_dir = r'./'
img_size=(300,300)
train_df=trim(train_df, max_samples, min_samples, column)
```

Original Number of classes in dataframe: 2
[3050, 3050]

```
In [19]: channels=3
batch_size=40
img_shape=(img_size[0], img_size[1], channels)
length=len(test_df)
test_batch_size=sorted([int(length/n) for n in range(1,length+1) if length %
test_steps=int(length/test_batch_size)
print ( 'test batch size: ',test_batch_size, ' test steps: ', test_steps)
def scalar(img):
    return img # EfficientNet expects pixels in range 0 to 255 so no scaling
trgen=ImageDataGenerator(preprocessing_function=scalar, horizontal_flip=True)
tvgen=ImageDataGenerator(preprocessing_function=scalar)
msg=' for the tr
print(msg, '\n', end='')
train_gen=trgen.flow_from_dataframe( train_df, x_col='filepaths', y_col='labels',
color_mode='rgb', shuffle=True, batch_size=40)
msg=' for the te
print(msg, '\n', end='')
test_gen=tvgen.flow_from_dataframe( test_df, x_col='filepaths', y_col='labels',
color_mode='rgb', shuffle=False, batch_size=40)
msg=' for the val
print(msg, '\n', end='')
valid_gen=tvgen.flow_from_dataframe( valid_df, x_col='filepaths', y_col='labels',
color_mode='rgb', shuffle=True, batch_size=40)
classes=list(train_gen.class_indices.keys())
class_count=len(classes)
train_steps=int(np.ceil(len(train_gen.labels)/batch_size))
labels=test_gen.labels
```

test batch size: 6 test steps: 89
Found 6100 validated image filenames belonging to 2 classes. for the train generator
Found 534 validated image filenames belonging to 2 classes. for the test generator
Found 533 validated image filenames belonging to 2 classes. for the validation generator

```
In [20]: ▶ show_image_samples(train_gen)
```




```
In [21]: ▶ model_name='EfficientNetB3'
base_model=tf.keras.applications.efficientnet.EfficientNetB3(include_top=False)
x=base_model.output
x=keras.layers.BatchNormalization(axis=-1, momentum=0.99, epsilon=0.001)(x)
x = Dense(256, kernel_regularizer = regularizers.l2(l = 0.016), activity_regularizer=regularizers.l1(0.006), activation='relu')(x)
x=Dropout(rate=.45, seed=123)(x)
output=Dense(class_count, activation='softmax')(x)
model=Model(inputs=base_model.input, outputs=output)
model.compile(Adamax(learning_rate=.001), loss='categorical_crossentropy', me
```

```
In [22]: ▶ epochs =40
patience= 1 # number of epochs to wait to adjust lr if monitored value does not improve
stop_patience =3 # number of epochs to wait before stopping training if monitored value does not improve
threshold=.9 # if train accuracy is < threshold adjust monitor accuracy, else use train accuracy
factor=.5 # factor to reduce lr by
dwell=True # experimental, if True and monitored metric does not improve on cross-validation, freeze weights
freeze=False # if true freeze weights of the base model
ask_epoch=5 # number of epochs to run before asking if you want to halt training
batches=train_steps
callbacks=[LRA(model=model, base_model= base_model, patience=patience, stop_patience=stop_patience, factor=factor, dwell=dwell, batches=batches, initial_epoch=0)]
```

```
In [23]: history=model.fit(x=train_gen, epochs=epochs, verbose=0, callbacks=callbacks,
                           validation_steps=None, shuffle=False, initial_epoch=0)
```

initializing callback starting training with base_model trainable

Epoch	Loss	Accuracy	V_loss	V_acc	LR	Next LR	Monitor	%
Improv	Duration							
1 /40	5.482	79.246	4.07478	74.484	0.00100	0.00100	accuracy	0.00
2 /40	2.834	87.311	2.53896	74.296	0.00100	0.00100	accuracy	10.18
3 /40	1.747	90.082	2.21746	68.668	0.00100	0.00100	val_loss	12.66
4 /40	1.138	91.902	0.99870	87.992	0.00100	0.00100	val_loss	54.96
5 /40	0.775	93.590	0.82506	82.364	0.00100	0.00100	val_loss	17.39

enter H to halt training or an integer for number of epochs to run then ask again


H

training has been halted at epoch 5 due to user input

Training is completed - model is set with weights from epoch 5

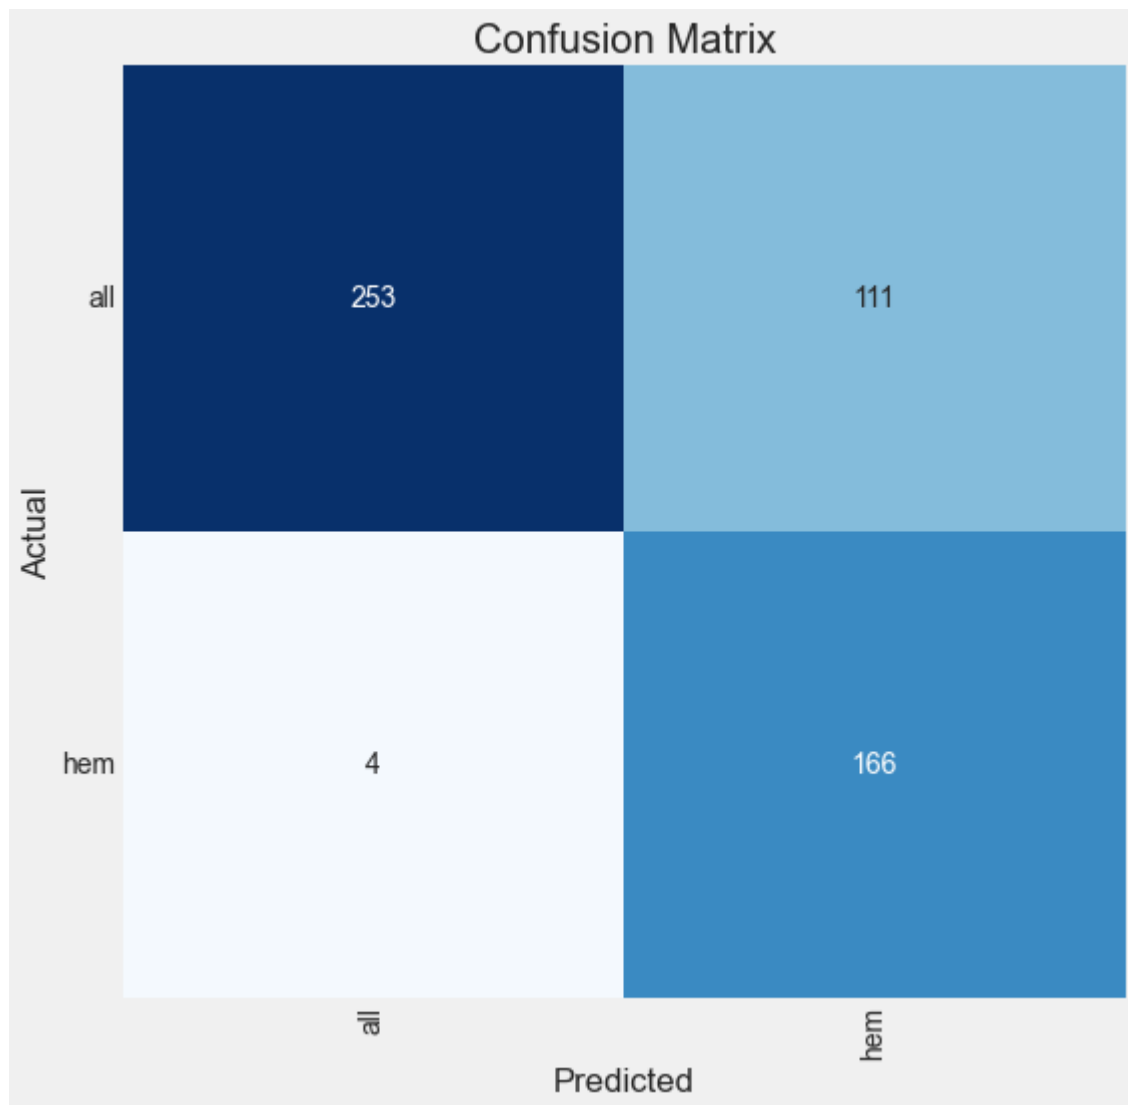

training elapsed time was 41.0 hours, 53.0 minutes, 49.35 seconds)

```
In [24]: tr_plot(history,0)
```



```
In [25]: ▶ subject='leukemia'
print_code=0
preds=model.predict(test_gen)
acc=print_info( test_gen, preds, print_code, working_dir, subject )
```

There were 115 errors in 534 test cases Model accuracy= 78.46 %



Classification Report:

```
-----
              precision    recall  f1-score   support

    all       0.9844      0.6951      0.8148        364
```

hem	0.5993	0.9765	0.7427	170
accuracy			0.7846	534
macro avg	0.7919	0.8358	0.7788	534
weighted avg	0.8618	0.7846	0.7919	534


```
In [26]: ▶ model_save_loc, csv_save_loc=saver(working_dir, model, model_name, subject, a
model was saved as ./EfficientNetB3-leukemia-0.78.h5
class csv file was saved as ./class_dict.csv
```

```

In [27]: ▶ img=plt.imread(img_path)
print ('Input image shape is ', img.shape)
# resize the image so it is the same size as the images the model was trained on
img=cv2.resize(img, img_size) # in earlier code img_size=(224,224) was used for
print ('the resized image has shape ', img.shape)
### show the resized image
plt.axis('off')
plt.imshow(img)
# Normally the next line of code rescales the images. However the EfficientNet
# img= img/255
# plt.imread returns a numpy array so it is not necessary to convert the image
# since we have only one image we have to expand the dimensions of img so it
# where the first dimension 1 is the batch size used by model.predict
img=np.expand_dims(img, axis=0)
print ('image shape after expanding dimensions is ',img.shape)
# now predict the image
pred=model.predict(img)
print ('the shape of prediction is ', pred.shape)
# this dataset has 15 classes so model.predict will return a list of 15 probabilities
# we want to find the index of the column that has the highest probability
index=np.argmax(pred[0])
# to get the actual Name of the class earlier I made a list of the class names
class=classes[index]
# Lets get the value of the highest probability
probability=pred[0][index]*100
# print out the class, and the probability
print(f'the image is predicted as being {class} with a probability of {probability}%')

```

Input image shape is (450, 450, 3)
 the resized image has shape (300, 300, 3)
 image shape after expanding dimensions is (1, 300, 300, 3)
 the shape of prediction is (1, 2)
 the image is predicted as being all with a probability of 90.25 %



```

In [28]: testdir=r'C:/Users/mahir/Documents/All the dreams in the world/Thesis Project
filepaths=[]
klass=[]
flist=os.listdir(testdir)
print ('number of test files is ', len(flist))
for f in flist:
    fpath=os.path.join(testdir,f)
    filepaths.append(fpath)
    klass=' '
Fseries=pd.Series(filepaths, name='filepaths')
Lseries=pd.Series(klass, name='Class')
test_df=pd.concat([Fseries, Lseries], axis=1)
length=len(test_df)
test_batch_size=sorted([int(length/n) for n in range(1,length+1) if length %
test_steps=int(length/test_batch_size)
print ( 'test batch size: ', test_batch_size, ' test steps: ', test_steps)
test_gen=tvgen.flow_from_dataframe( test_df, x_col='filepaths', y_col=None, t
                                color_mode='rgb', shuffle=False, batch_si
preds=model.predict(test_gen, verbose=1, steps=test_steps)
for i, p in enumerate(preds):
    index=np.argmax(p)
    klass=classes[index]
    test_df['Class'].iloc[i]=klass
print (test_df.head())
csv_name='submission'
csv_path=os.path.join(working_dir, csv_name)
test_df.to_csv(csv_path, index=False)
# read in csv file to see if it is correct
df_test=pd.read_csv(csv_path)
print (df_test.head())

```

```

number of test files is 2586
test batch size: 6 test steps: 431
Found 2586 validated image filenames.
431/431 [=====] - 760s 2s/step
filepaths Class
0 C:/Users/mahir/Documents/All the dreams in the... hem
1 C:/Users/mahir/Documents/All the dreams in the... hem
2 C:/Users/mahir/Documents/All the dreams in the... hem
3 C:/Users/mahir/Documents/All the dreams in the... hem
4 C:/Users/mahir/Documents/All the dreams in the... hem
filepaths Class
0 C:/Users/mahir/Documents/All the dreams in the... hem
1 C:/Users/mahir/Documents/All the dreams in the... hem
2 C:/Users/mahir/Documents/All the dreams in the... hem
3 C:/Users/mahir/Documents/All the dreams in the... hem
4 C:/Users/mahir/Documents/All the dreams in the... hem

```

In []: 

